

Google DeepMind



Attention Sink in LLMs and its Applications

Presenter: **Xiangming Gu**

Affiliation: Google Deepmind, National University of Singapore

I am attempting to answer ...

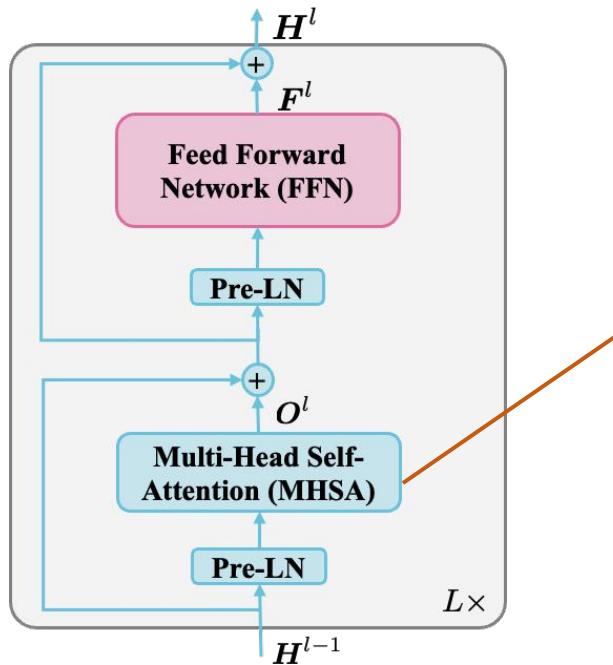
- Mechanism understanding of Attention Sink?
- When Attention Sink Emerges in LLMs?
- Why LLMs need Attention Sink?
- Why GPT-OSS and Qwen3-Next consider Attention Sink in the Model Design?

Covered the following two papers

- When Attention Sink Emerges in Language Models: An Empirical View. ICLR 2025
- Why Do LLMs Attend to the First Token? COLM 2025

What is Attention Sink?

- Decoder-only Transformer



Self-attention is one of the most important parts

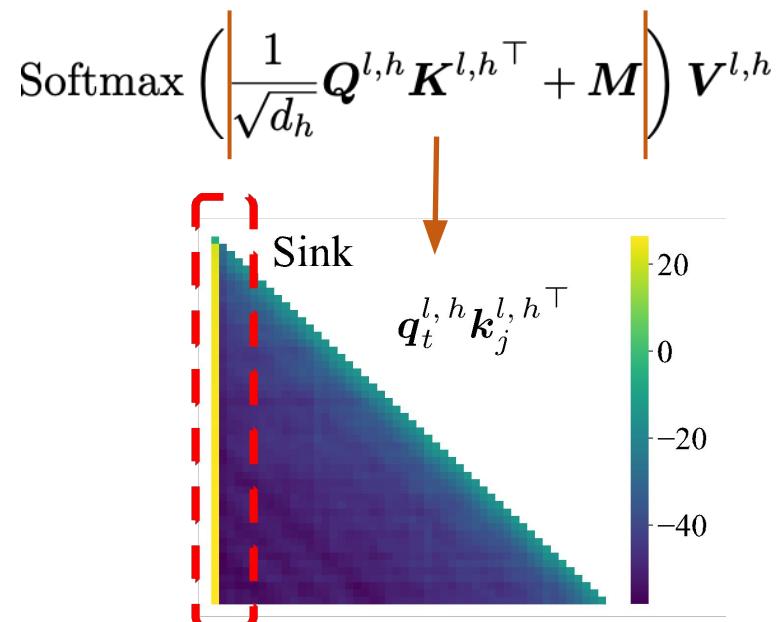
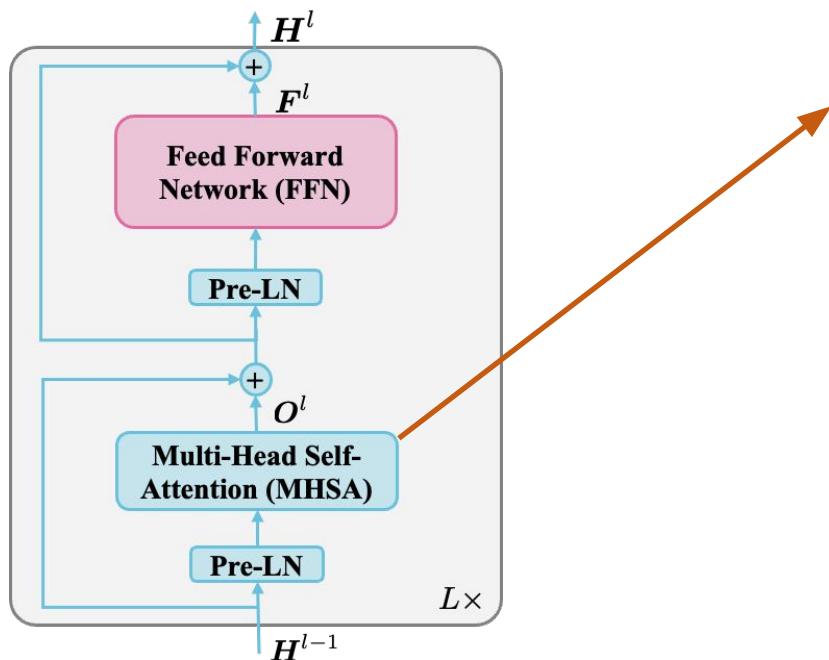
$$\text{Softmax}\left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \mathbf{K}^{l,h\top} + \mathbf{M}\right) \mathbf{V}^{l,h}$$



queries keys values

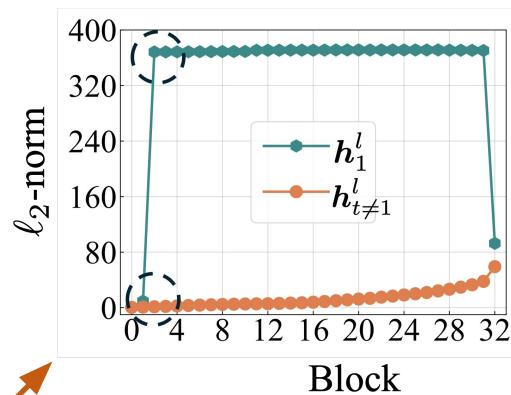
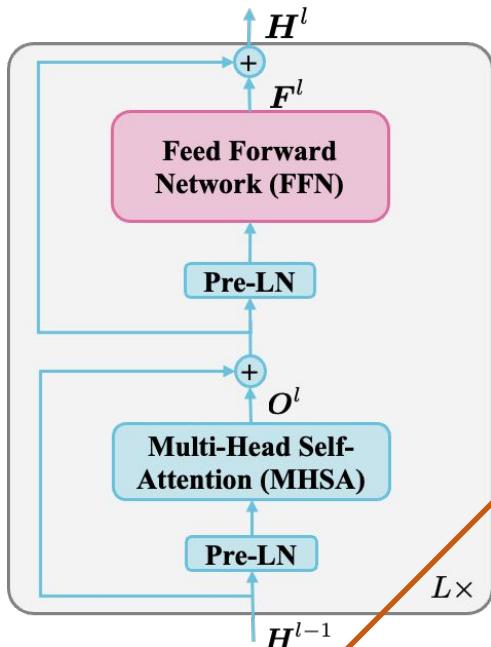
Casual mask

What is Attention Sink?

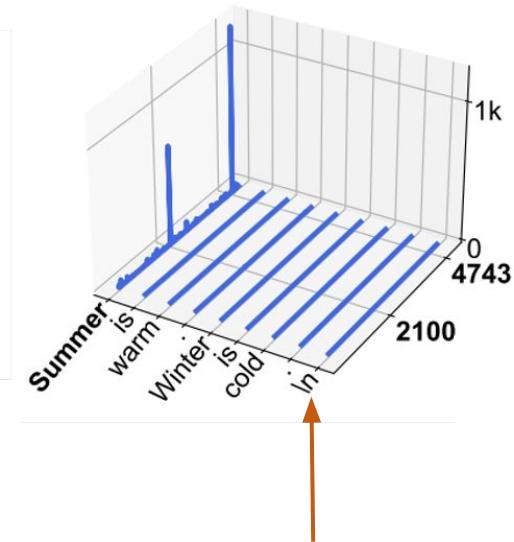


Phenomenons associated to Attention Sink

- Massive Activations



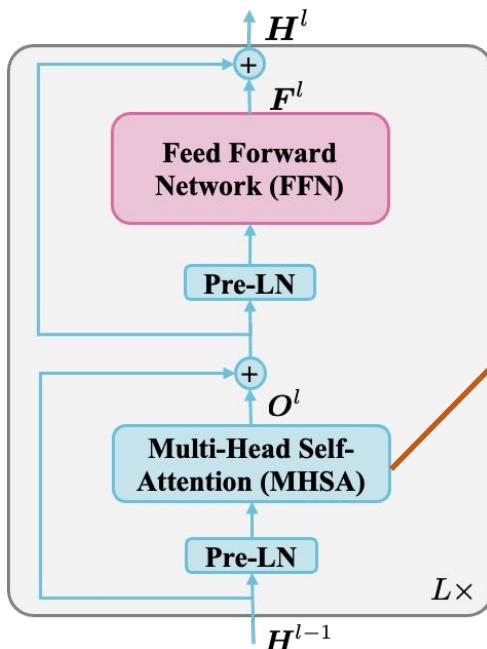
Activations extremely large



Few dimensions have spikes/outliers

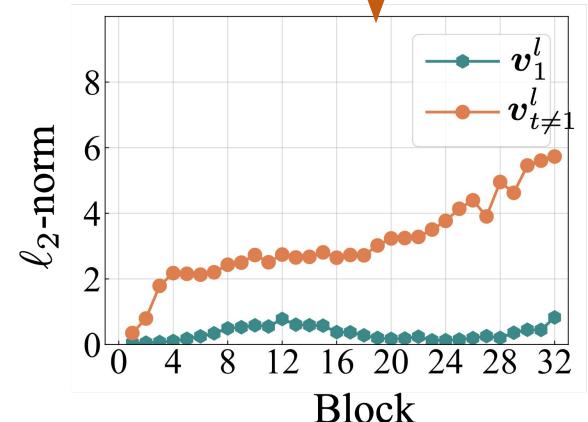
Phenomenons associated to Attention Sink

- Value Drains



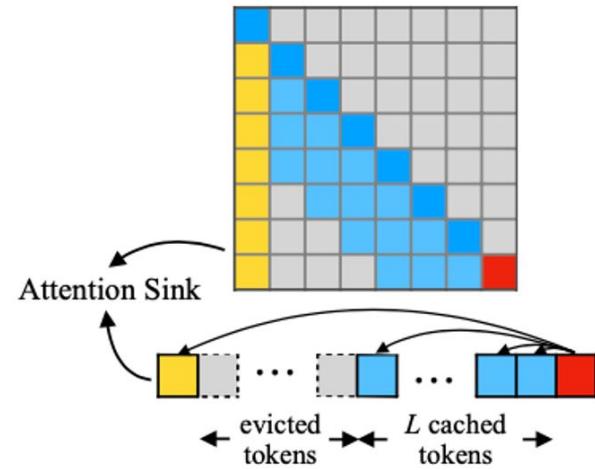
$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \mathbf{K}^{l,h \top} + \mathbf{M} \right) \mathbf{V}^{l,h}$$

Values extremely small



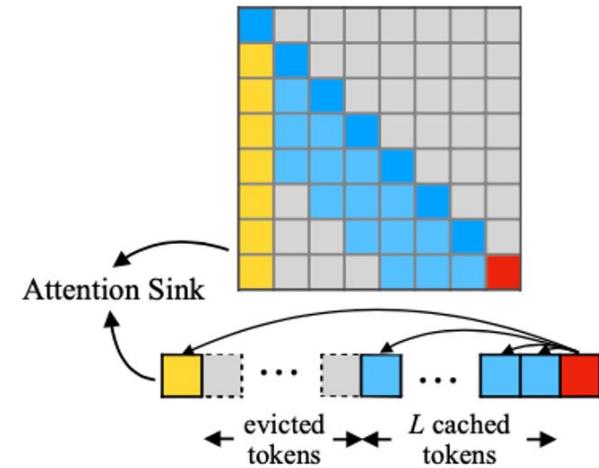
Post-hoc Applications of Attention Sink

- Long context understanding / generation
- Only computing attention on the first token and recent tokens



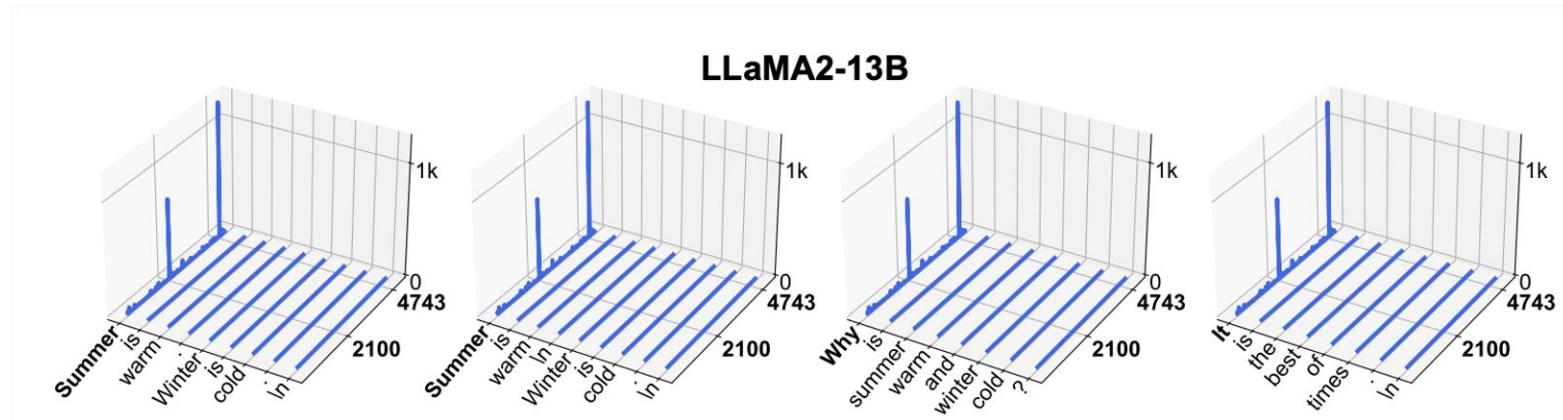
Post-hoc Applications of Attention Sink

- KV cache optimization
- Only retaining KV cache of sink tokens and recent tokens



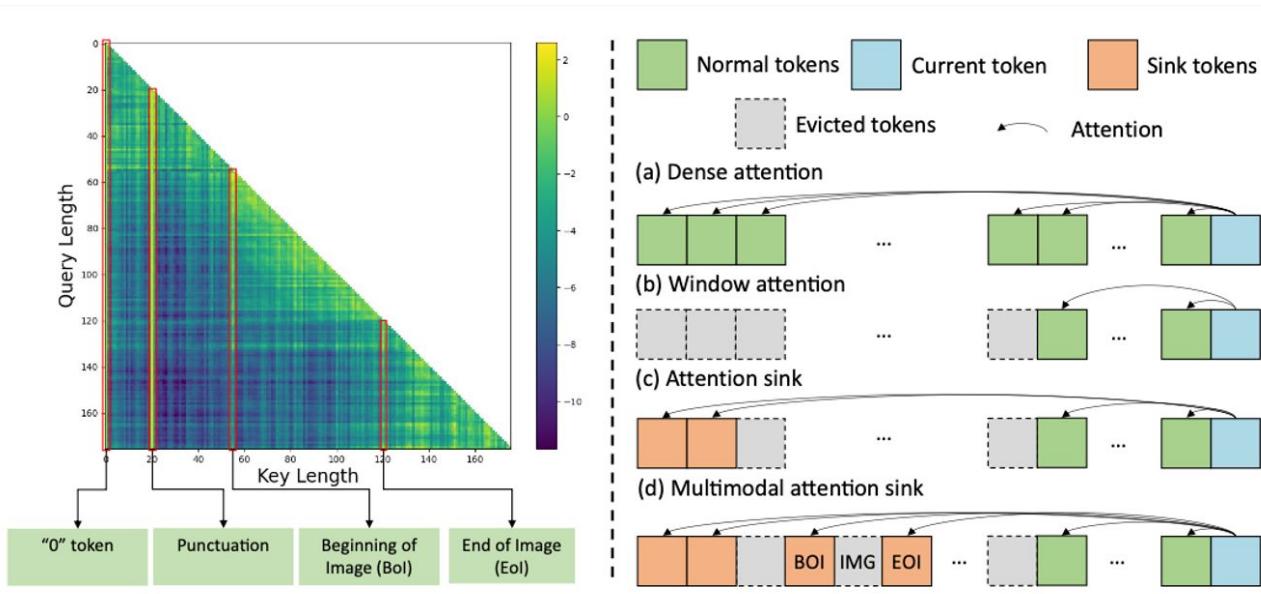
Post-hoc Applications of Attention Sink

- Model quantization
- Preserving the full precision of KV cache of sink token



Post-hoc Applications of Attention Sink

- Multimodal language modeling



I am attempting to answer ...

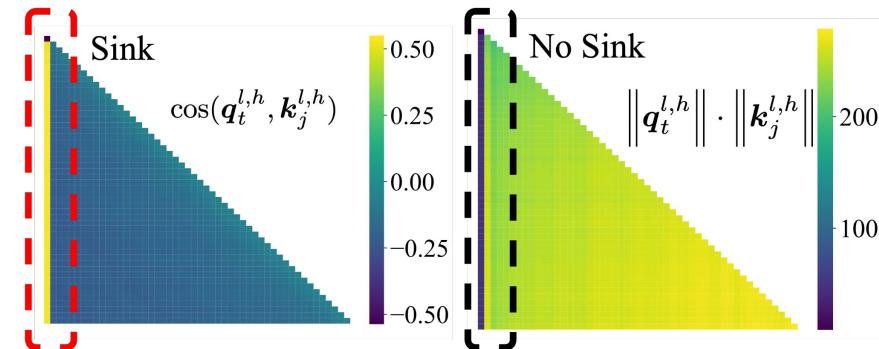
- Mechanism understanding of Attention Sink?
- When Attention Sink Emerges in LLMs?
- Why LLMs need Attention Sink?
- Why GPT-OSS and Qwen3-Next consider Attention Sink in the Model Design?

Mechanism Understanding of Attention Sink

Attention sink is due to the key **key** bias of the sink token

$$\mathbf{q}_t^{l,h} \mathbf{k}_1^{l,h}^\top \gg \mathbf{q}_t^{l,h} \mathbf{k}_{j \neq 1}^{l,h}^\top$$

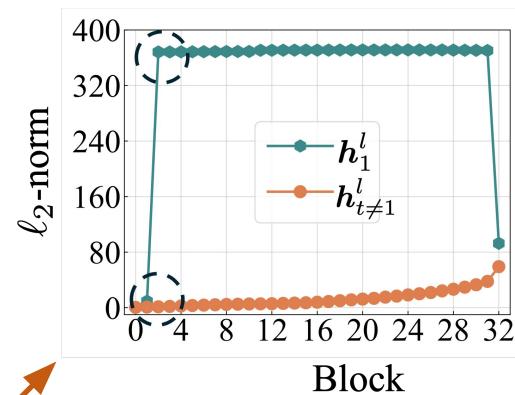
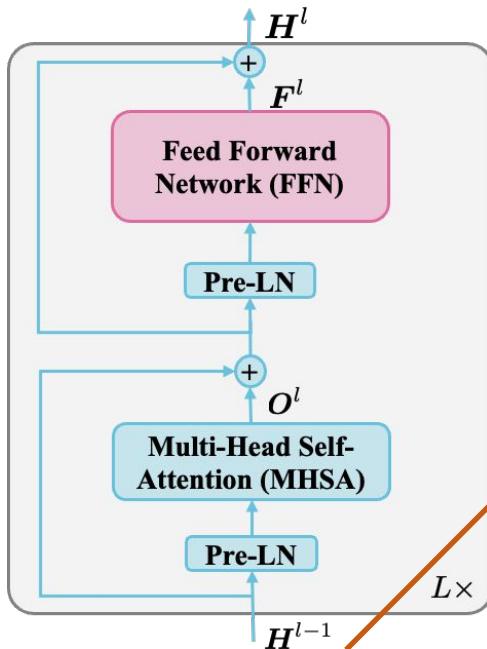
$$\cos(\mathbf{q}_t^{l,h}, \mathbf{k}_1^{l,h}) \gg \cos(\mathbf{q}_t^{l,h}, \mathbf{k}_{j \neq 1}^{l,h})$$



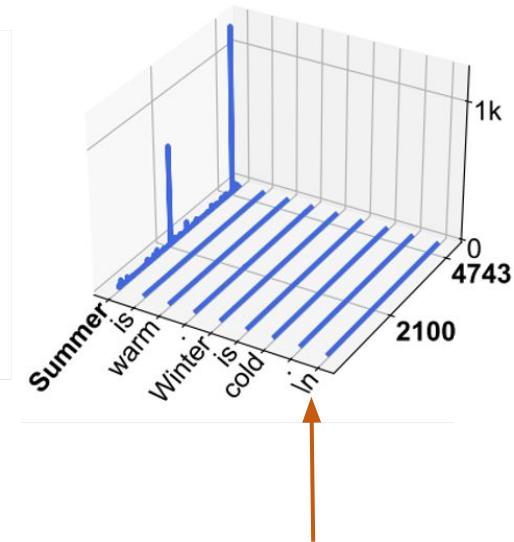
key of the sink token is located in the different manifold, it has small angles with any queries

Mechanism Understanding of Attention Sink

- Massive Activations



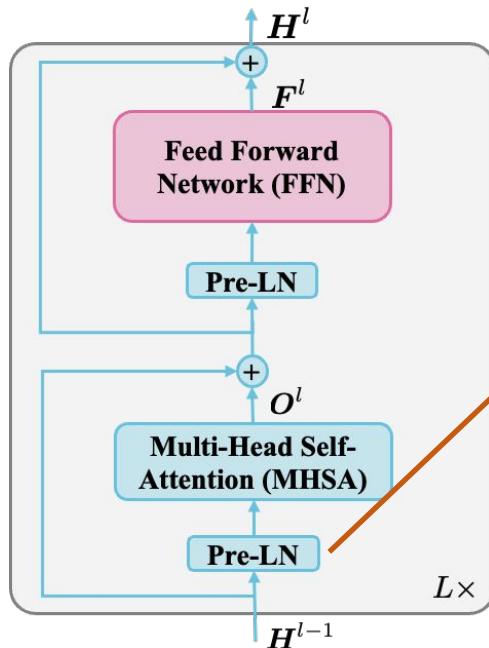
Activations extremely large



Few dimensions have spikes/outliers

Mechanism Understanding of Attention Sink

- Existence of massive activations is to support attention sink



$$\text{LN}(h) = \frac{h}{\sqrt{\frac{1}{d} \sum_{i=1}^d h_i^2}} \odot g$$

Layer Norm only retain the spike dimensions (dominate the norm)

$$k_t^{l,h} = \text{LN}(h_t^{l-1}) W_K^{l,h} R_{\Theta, -t}$$

Linear transformations of spikes

Similar mechanism for small values

Mechanism Understanding of Attention Sink

Why all these phenomenon tend to happen in the first token (not necessary to be BOS)?

- **Uniqueness of the first token:** self-attention involves no other tokens, all hidden states in the forward path are equivalent to MLP transformations of input embeddings
- LLMs learn to map the input embeddings to massive activations after certain layers, leading to key bias, and then attention sink

Mechanism Understanding of Attention Sink

- Attention sink approximates “no-op”

$$\mathbf{v}_i^\dagger = \sum_{j=1}^i \alpha_{ij} \mathbf{v}_j = \alpha_{i1} \boxed{\mathbf{v}_1} + \sum_{j \neq 1}^i \boxed{\alpha_{ij}} \mathbf{v}_j$$

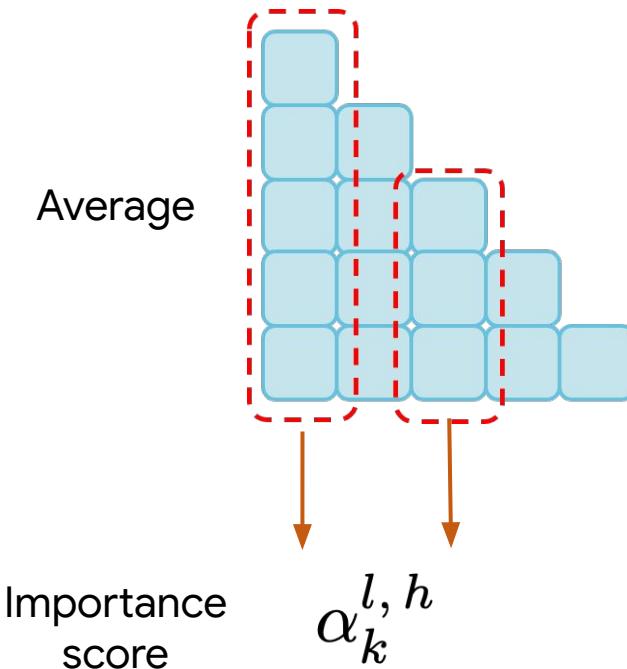
Small

I am attempting to answer ...

- Mechanism understanding of Attention Sink?
- When Attention Sink Emerges in LLMs?
- Why LLMs need Attention Sink?
- Why GPT-OSS and Qwen3-Next consider Attention Sink in the Model Design?

A metric to measure Attention Sink

- Motivations: attention scores of the first token dominates



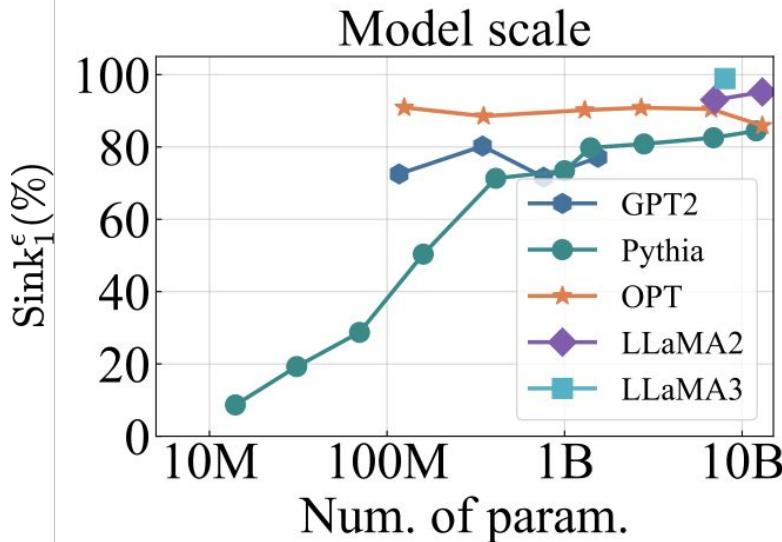
$$\text{Sink}_k^\epsilon = \frac{1}{L} \sum_{l=1}^L \frac{1}{H} \sum_{h=1}^H \mathbb{I}(\alpha_k^{l,h} > \epsilon)$$

Attention sink metric of the whole LM

Within a head, a threshold to decide sink, e.g., 0.3 for 64 tokens

Attention Sink w.r.t. Model Scale / Training Stage

- Attention sink emerges in small LMs, even with 14M params.
- Attention sink already emerges in LM pre-training.



LLM	Sink $_{\text{1}}^{\epsilon}$ (%)	
	Base	Chat
Mistral-7B	97.49	88.34
LLaMA2-7B	92.47	92.88
LLaMA2-13B	91.69	90.94
LLaMA3-8B	99.02	98.85

Attention Sink w.r.t. Different Inputs

- Attention sink emerges with / without BOS (for most LLMs), even with random tokens as input
- Under all the repeated tokens?

LLM	natural	$\text{Sink}_1^\epsilon(\%)$ random	repeat
GPT2-XL	77.00	70.29	62.28
Mistral-7B	97.49	75.21	0.00
LLaMA2-7B Base	92.47	90.13	0.00
LLaMA3-8B Base	99.02	91.23	0.00

Related to positional embeddings

Attention Sink with Repeated Tokens as Inputs

- For LLMs with NOPE / Relative PE / ALiBi / Rotary

$$\mathbf{P} = \mathbf{0}$$

Residual streams before Transformer blocks

$$\mathbf{h}_t^0 = \mathbf{xW}_E + \mathbf{P}$$

Then

$$\mathbf{h}_1^0 = \mathbf{h}_2^0 = \cdots = \mathbf{h}_T^0$$

Using induction, we can prove (all have massive activations, distribute the sink)

$$\mathbf{h}_1^l = \mathbf{h}_2^l = \cdots = \mathbf{h}_T^l, \quad \forall \ 0 \leq l \leq L$$

Attention Sink with Repeated Tokens as Inputs

- We can even derive the closed form / upper bound attention distributions for NOPE / Relative PE / ALiBi / Rotary (see the paper).
- However, absolute / learnable PE (e.g., GPT2) have no such properties

LLM	Sink ₁ ^ε (%)		
	natural	random	repeat
GPT2-XL	77.00	70.29	62.28
Mistral-7B	97.49	75.21	0.00
LLaMA2-7B Base	92.47	90.13	0.00
LLaMA3-8B Base	99.02	91.23	0.00

When Attention Sink Emerges in LLMs?

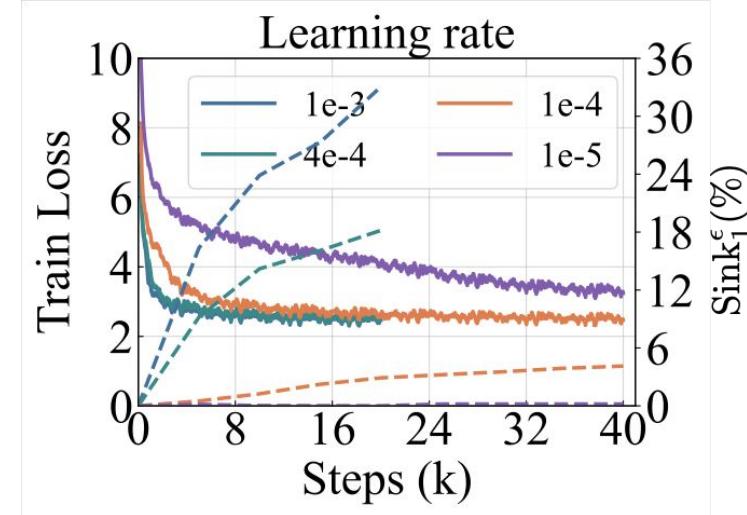
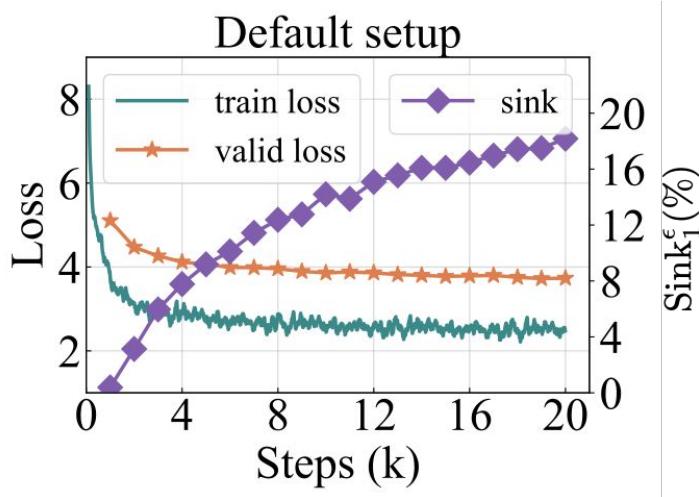
- Attention sink appears during LLM pre-training
- Attributing attention sink phenomenon to LLM pre-training

$$\min_{\theta} \mathbb{E}_{\mathbf{X} \sim p_{\text{data}}} [\mathcal{L}(p_{\theta}(\mathbf{X}))]$$

Optimization Data distribution Loss function Model architecture

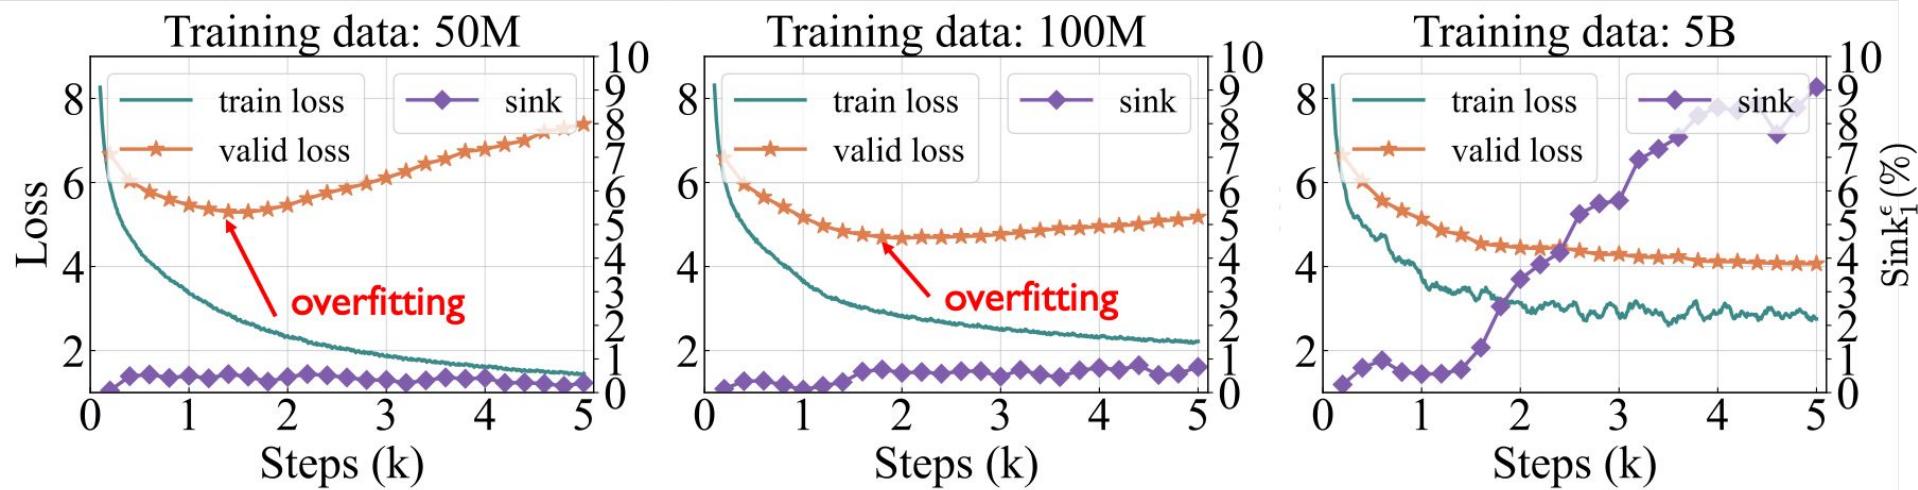
Effects of Optimization

- Attention sink appears during LLM pre-training process (not initialization)
- Large LR encourages attention sink (even under the same LR*steps)



Effects of Data Distribution

- Attention sink emerges when we have enough unique training data amount



Effects of Loss function

- Weight decay encourages attention sink

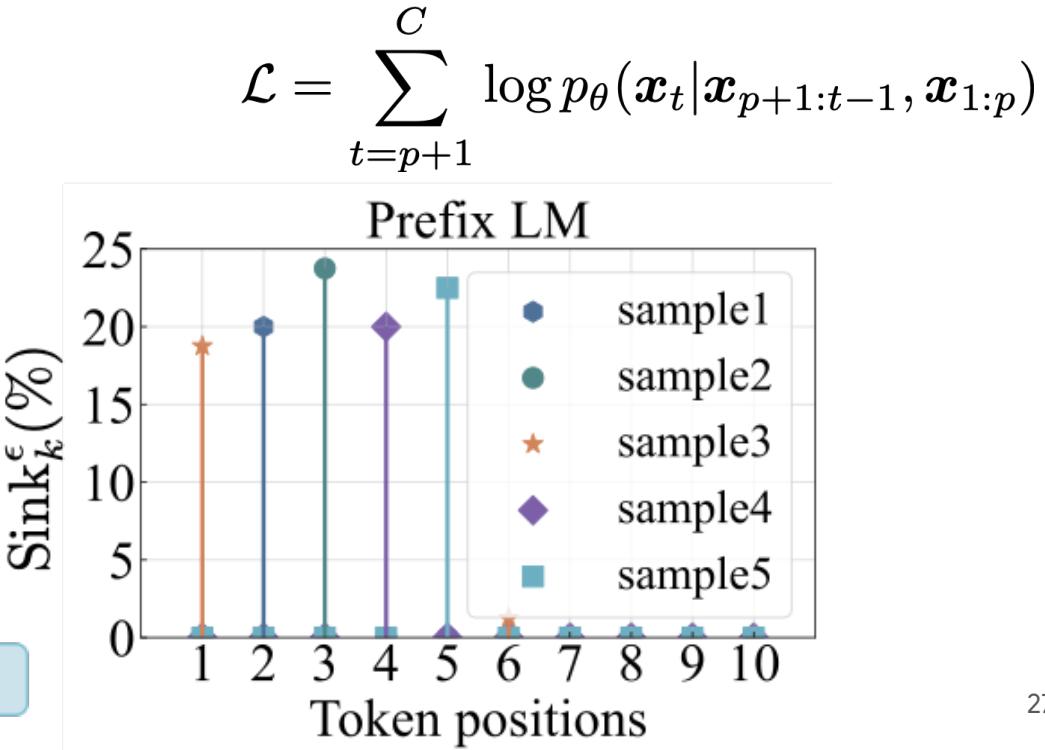
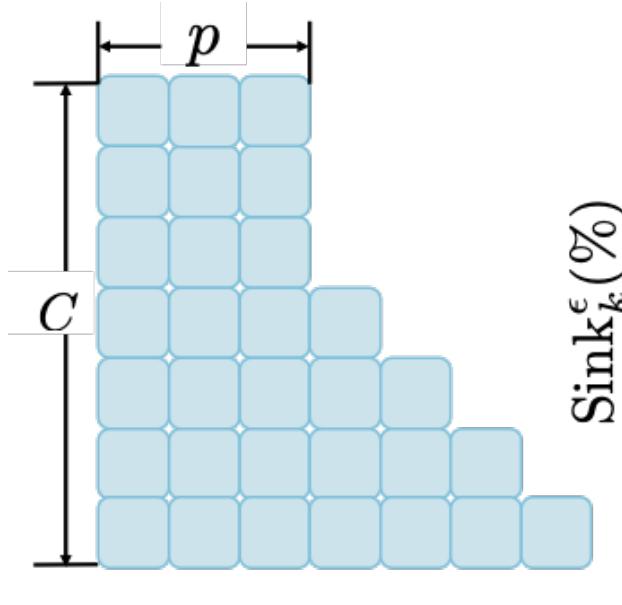
$$\mathcal{L} = \sum_{t=2}^C \log p_\theta(\mathbf{x}_t | \mathbf{x}_{<t}) + \gamma \|\theta\|_2^2$$

L2 regularization

γ	0.0	0.001	0.01	0.1	0.5	1.0	2.0	5.0
Sink $^\epsilon_1$ (%)	15.20	15.39	15.23	18.18	41.08	37.71	6.13	0.01
valid loss	3.72	3.72	3.72	3.73	3.80	3.90	4.23	5.24

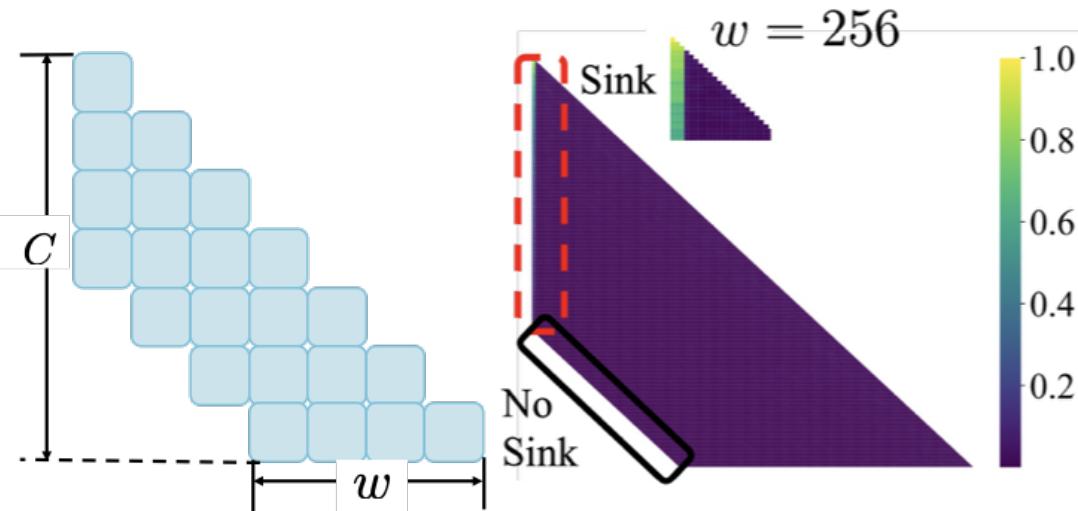
Effects of Loss function

- **Prefix language modeling:** sink token shifts from the first token to other positions within the prefix



Effects of Loss function

- Shift window attention: attention sink appears on the **absolute, not the relative first token**
- Small window size mitigates attention sink



Validating sink token has
key bias

$$\mathcal{L} = \sum_{t=2}^C \log p_{\theta}(\mathbf{x}_t | \mathbf{x}_{t-w:t-1})$$

Effects of Model Architecture

The following designs do not affect the emergence of attention sink

- Positional embeddings

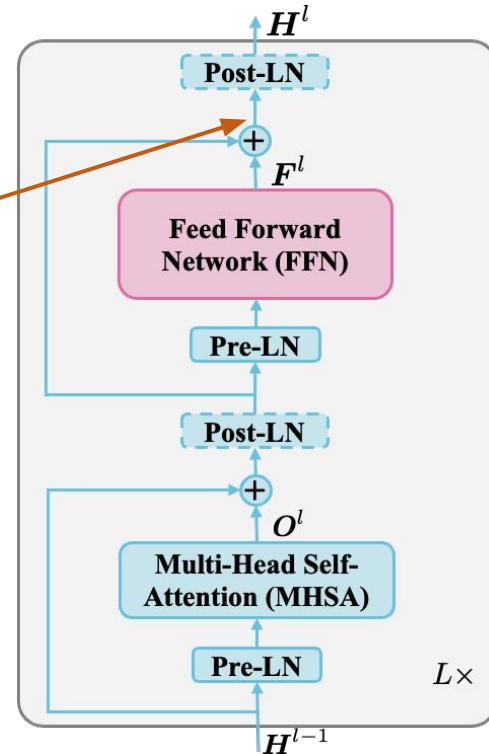
NOPE, learnable PE, absolute PE, relative PE, Rotary, ALIBI

Effects of Model Architecture

The following designs do not affect the emergence of attention sink

- Positional embeddings
- Pre-norm or post-norm

Massive activations
happen before LN



Effects of Model Architecture

The following designs do not affect the emergence of attention sink

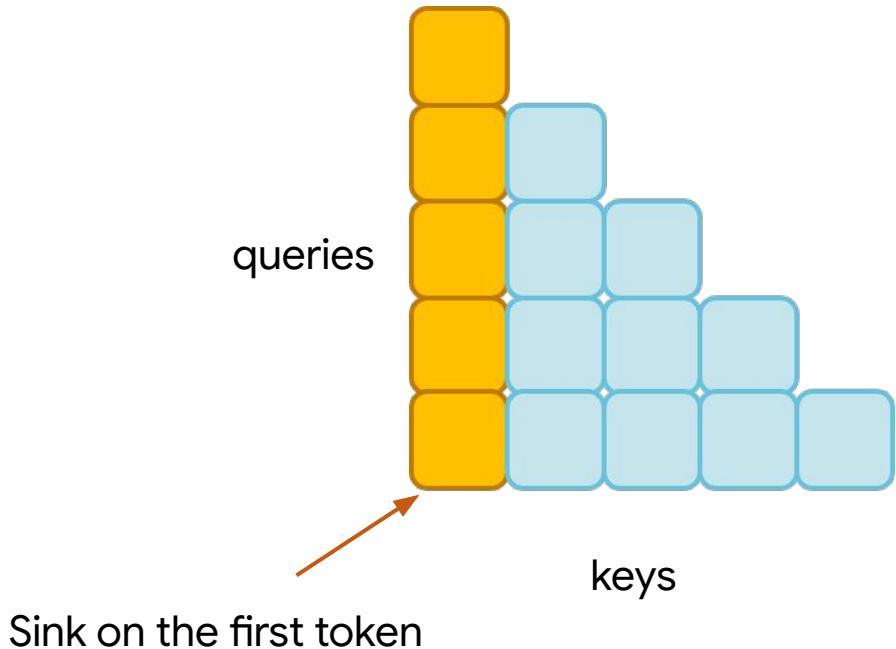
- Positional embeddings
- Pre-norm or post-norm
- FFNs with different activation functions
- Number of attention heads, how to combine multiple heads
- ...

Effects of Model Attention Design

- Standard softmax attention

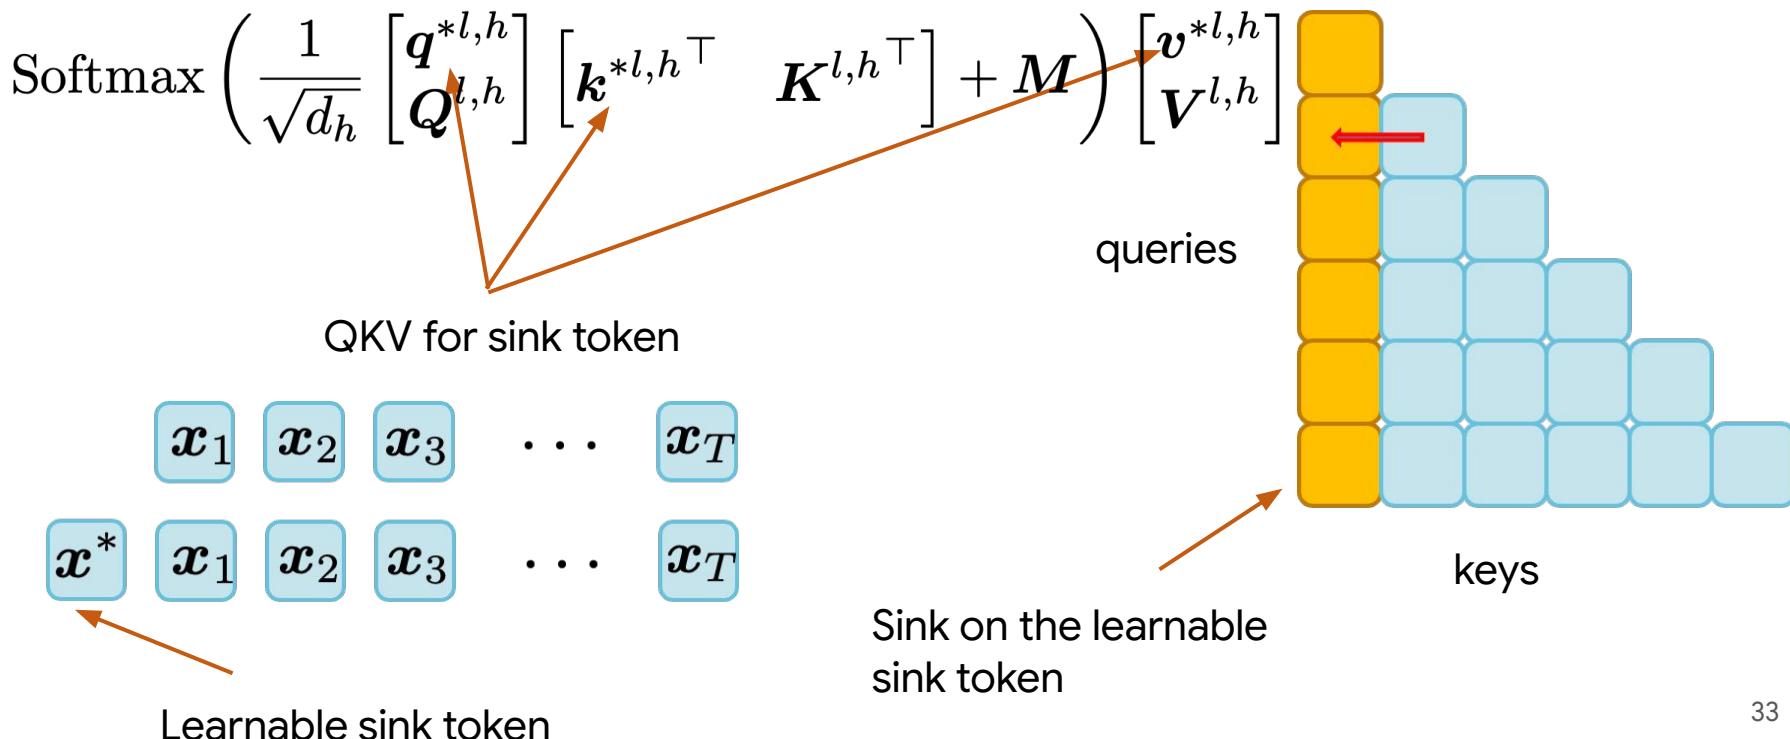
$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \mathbf{K}^{l,h^\top} + \mathbf{M} \right) \mathbf{V}^{l,h}$$

queries keys values
Casual mask



Effects of Model Attention Design

- Softmax attention with a **learnable sink token**



Effects of Model Attention Design

- Softmax attention with **learnable KV biases**

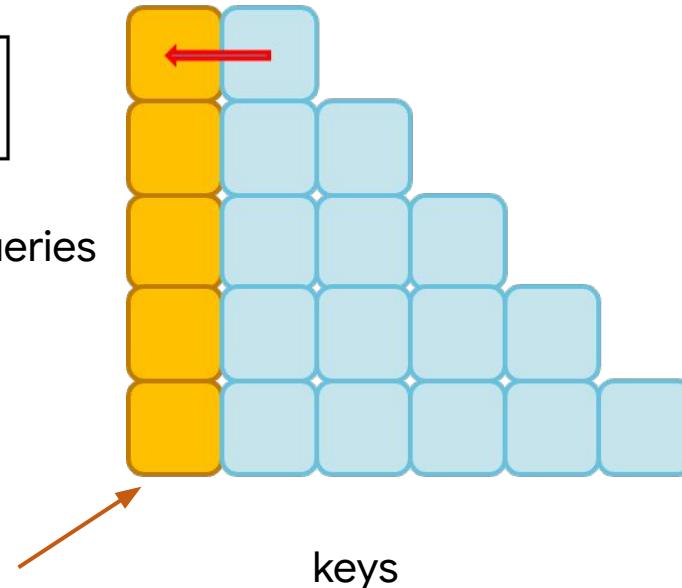
$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} Q^{l,h} \begin{bmatrix} k^{*l,h^\top} & K^{l,h^\top} \end{bmatrix} + M \right) \begin{bmatrix} v^{*l,h} \\ V^{l,h} \end{bmatrix}$$

Learnable KV biases

queries

keys

Sink on the learnable K
biases

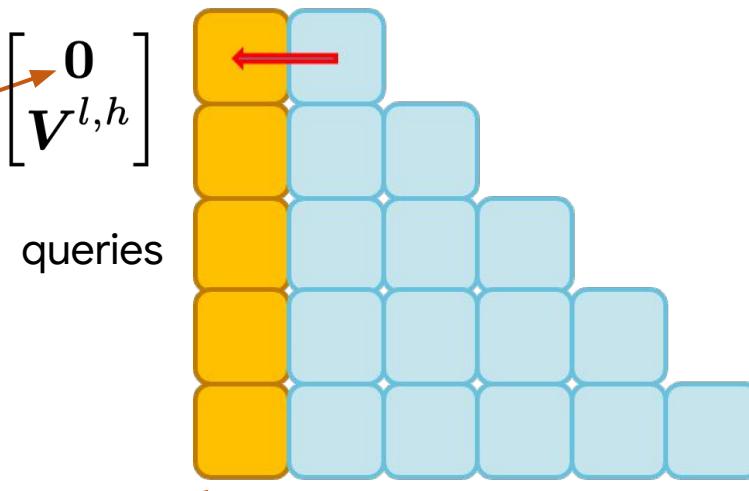


Effects of Model Attention Design

- Softmax attention with **learnable K biases**

$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} Q^{l,h} \begin{bmatrix} k^{*l,h^\top} & K^{l,h^\top} \end{bmatrix} + M \right) \begin{bmatrix} \mathbf{0} \\ V^{l,h} \end{bmatrix}$$

Learnable K biases,
zero V biases



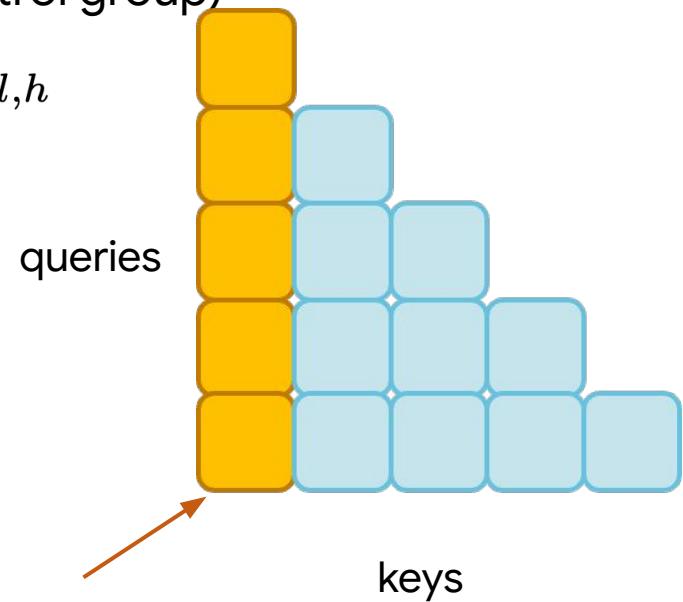
Sink on the learnable
key biases

Effects of Model Attention Design

- Softmax attention with **learnable K biases** (control group)

$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} Q^{l,h} K^{l,h^\top} + M \right) V^{l,h} + v^{*l,h}$$

Learnable V biases



Sink on the first token,
no effects

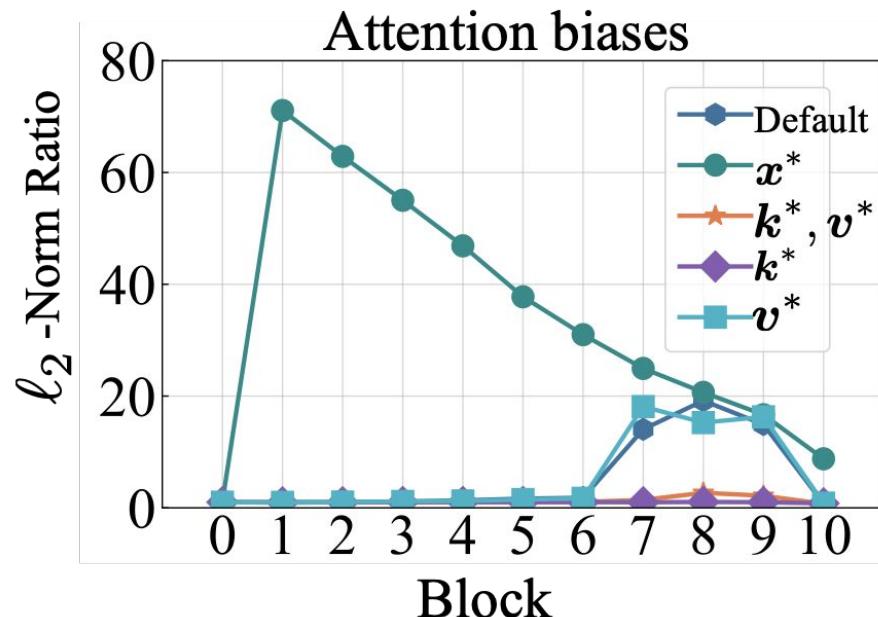
Effects of Attention Biases

- Attention biases can absorb attention sink from the actual first token

Attention in each head	Sink _* ^ε (%)	Sink ₁ ^ε (%)	valid loss
Softmax $\left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \mathbf{K}^{l,h \top} + \mathbf{M} \right) \mathbf{V}^{l,h}$	-	18.18	3.73
Softmax $\left(\frac{1}{\sqrt{d_h}} \begin{bmatrix} \mathbf{q}^{*l,h} \\ \mathbf{Q}^{l,h} \end{bmatrix} \begin{bmatrix} \mathbf{k}^{*l,h \top} & \mathbf{K}^{l,h \top} \end{bmatrix} + \mathbf{M} \right) \begin{bmatrix} \mathbf{v}^{*l,h} \\ \mathbf{V}^{l,h} \end{bmatrix}$	74.12	0.00	3.72
Softmax $\left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \begin{bmatrix} \mathbf{k}^{*l,h \top} & \mathbf{K}^{l,h \top} \end{bmatrix} + \mathbf{M} \right) \begin{bmatrix} \mathbf{v}^{*l,h} \\ \mathbf{V}^{l,h} \end{bmatrix}$	72.76	0.04	3.72
Softmax $\left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \begin{bmatrix} \mathbf{k}^{*l,h \top} & \mathbf{K}^{l,h \top} \end{bmatrix} + \mathbf{M} \right) \begin{bmatrix} \mathbf{0} \\ \mathbf{V}^{l,h} \end{bmatrix}$	73.34	0.00	3.72
Softmax $\left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \mathbf{K}^{l,h \top} + \mathbf{M} \right) \mathbf{V}^{l,h} + \mathbf{v}^{*l,h}$	-	17.53	3.73

Effects of Attention Biases

- Key biases can significantly mitigate massive activations, as no need to develop new biases



Effects of Attention Biases

- Value bias needs to be close to zero

$v^{*l,h}$	0	v'	$5v'$	$20v'$	v''	$5v''$	$20v''$
Sink $_{*}^{\epsilon}$ (%)	73.34	70.03	44.43	1.51	69.74	27.99	0.00
Sink $_{1}^{\epsilon}$ (%)	0.00	0.06	3.71	25.88	2.15	5.93	11.21
valid loss	3.72	3.72	3.72	3.71	3.72	3.72	3.73

$$v' = [1, 0, 0, \dots, 0]$$

$$v'' = [1, 1, 1, \dots, 1] / \sqrt{d_h}$$

Effects of Attention Biases

- Key bias is low-rank

d_a	1	2	4	8	16	32	64
$\text{Sink}_*^\epsilon(\%)$	32.18	30.88	30.94	31.39	23.30	51.23	69.19
$\text{Sink}_1^\epsilon(\%)$	4.74	4.96	4.39	4.54	2.19	1.94	0.04
valid loss	3.73	3.72	3.72	3.73	3.73	3.73	3.72

Comparing different Attention Biases

- Learnable key biases, zero value biases

$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \begin{bmatrix} \mathbf{k}^{*l,h^\top} & \mathbf{K}^{l,h^\top} \end{bmatrix} + \mathbf{M} \right) \begin{bmatrix} \mathbf{0} \\ \mathbf{V}^{l,h} \end{bmatrix}$$

- Softmax off-by-one

$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} \begin{bmatrix} \mathbf{0}^{*l,h^\top} & \mathbf{Q}^{l,h} \mathbf{K}^{l,h^\top} \end{bmatrix} + \mathbf{M} \right) \begin{bmatrix} \mathbf{0} \\ \mathbf{V}^{l,h} \end{bmatrix}$$

- Learnable attention score biases (single number for each head, layer)

$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} \begin{bmatrix} \mathbf{b}^{*l,h^\top} & \mathbf{Q}^{l,h} \mathbf{K}^{l,h^\top} \end{bmatrix} + \mathbf{M} \right) \begin{bmatrix} \mathbf{0} \\ \mathbf{V}^{l,h} \end{bmatrix}$$

$$\mathbf{b}^{*l,h} = b^{*l,h}[1, 1, 1, \dots, 1]$$

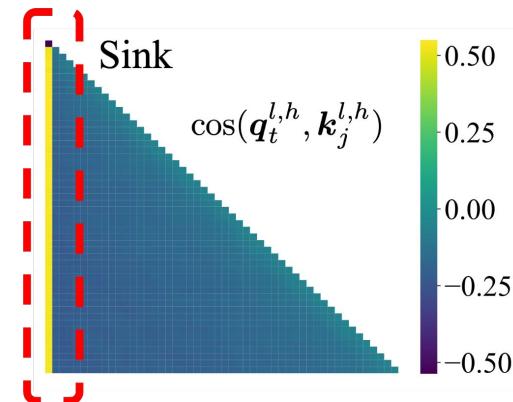
Comparing different Attention Biases

- Softmax off-by-one: with any query, the cosine similarity is zero

$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} \begin{bmatrix} \mathbf{0}^{*l,h^\top} & \mathbf{Q}^{l,h} \mathbf{K}^{l,h^\top} \end{bmatrix} + \mathbf{M} \right) \begin{bmatrix} \mathbf{0} \\ \mathbf{V}^{l,h} \end{bmatrix}$$

- Original format:
 - Zero may already be enough

$$(\text{softmax}_1(x))_i = \frac{\exp(x_i)}{1 + \sum_j \exp(x_j)}$$



Effects of Attention Biases

The learnable key bias and **zero** value bias experiments show that:

- Large attention score does not mean important in semantic
- Sink token save extra attention, adjusts the dependence among tokens

But why LLMs need such a mechanism?

Effects of Normalization in Softmax Attention

Whether this is due to the normalization in Softmax attention?

$$\begin{aligned}\mathbf{v}_i^\dagger &= \sum_{j=1}^i \frac{\alpha \text{sim}(\varphi(\mathbf{q}_i), \varphi(\mathbf{k}_j))}{\sum_{j'=1}^i \text{sim}(\varphi(\mathbf{q}_i), \varphi(\mathbf{k}_{j'}))} \mathbf{v}_j = \sum_{j=1}^i \frac{\text{sim}(\varphi(\mathbf{q}_i), \varphi(\mathbf{k}_j))}{\sum_{j'=1}^i \text{sim}(\varphi(\mathbf{q}_i), \varphi(\mathbf{k}_{j'}))} \mathbf{h}_j(\alpha \mathbf{W}_V), \\ \mathbf{o}'_i &= \text{Concat}_{h=1}^H (\mathbf{v}_i'^h) \mathbf{W}_O.\end{aligned}$$

Scaling the normalization $\mathbf{Z}_i \rightarrow \mathbf{Z}_i/\alpha$, equivalent to scaling weight matrices, and then scaling the LR, mitigates attention sink

$$\begin{aligned}\mathbf{W}_O^{s+1} &= \mathbf{W}_O^s - \eta \nabla_{\mathbf{W}_O^s} \mathcal{L}(\alpha \mathbf{W}_O^s) \\ &= \mathbf{W}_O^s - \alpha \eta \nabla_{\mathbf{W}} \mathcal{L}(\mathbf{W})|_{\mathbf{W}=\alpha \mathbf{W}_O^s},\end{aligned}$$

$$\begin{aligned}\hat{\mathbf{W}}_O^{s+1} &= \hat{\mathbf{W}}_O^s - \eta' \nabla_{\hat{\mathbf{W}}_O^s} \mathcal{L}(\hat{\mathbf{W}}_O^s) \\ &= \alpha \mathbf{W}_O^s - \eta' \nabla_{\mathbf{W}} \mathcal{L}(\mathbf{W})|_{\mathbf{W}=\alpha \mathbf{W}_O^s},\end{aligned}$$

Effects of Normalization in Softmax Attention

Power of sum to one: may mitigate attention sink but does not prevent, sensitive to LR, large LR may incentivize attention sink

$$\mathbf{v}_i^\dagger = \frac{\sum_{j=1}^i \text{sim}(\varphi(\mathbf{q}_i), \varphi(\mathbf{k}_j)) \mathbf{v}_j}{\left(\sum_{j'=1}^i \text{sim}(\varphi(\mathbf{q}_i), \varphi(\mathbf{k}_{j'}))^p \right)^{\frac{1}{p}}}$$
$$\mathbf{v}_i^\dagger = \sum_{j=1}^i \left(\frac{\exp\left(\frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d_h}/p}\right)}{\sum_{j'=1}^i \exp\left(\frac{\mathbf{q}_i \mathbf{k}_{j'}^\top}{\sqrt{d_h}/p}\right)} \right)^{\frac{1}{p}} \mathbf{v}_j$$

Effects of Normalization in Softmax Attention

- Removing the normalization in Softmax attention

Using sigmoid attention (exponential kernel in Softmax tends to explode)

$$\text{Sigmoid} \left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \mathbf{K}^{l,h \top} + \mathbf{M} \right) \mathbf{V}^{l,h}$$

Or ELU plus one attention

No normalization -> No attention sink; add back -> attention sink

Effects of Normalization in Softmax Attention

Other attention variants

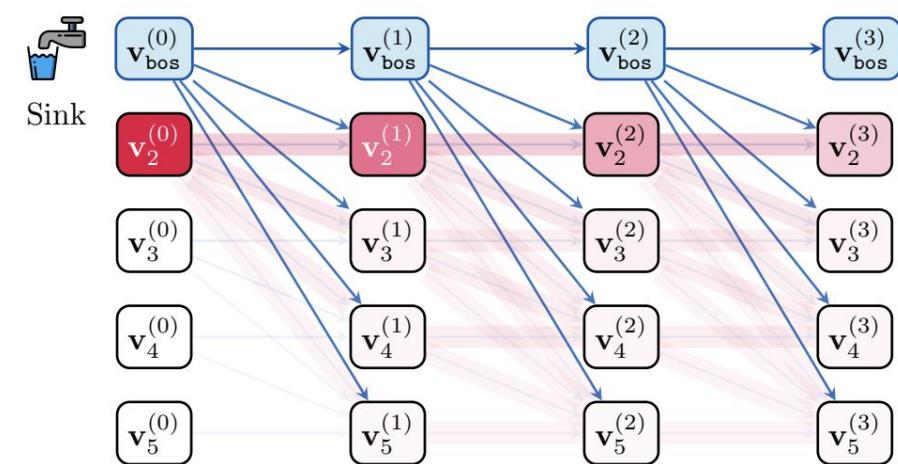
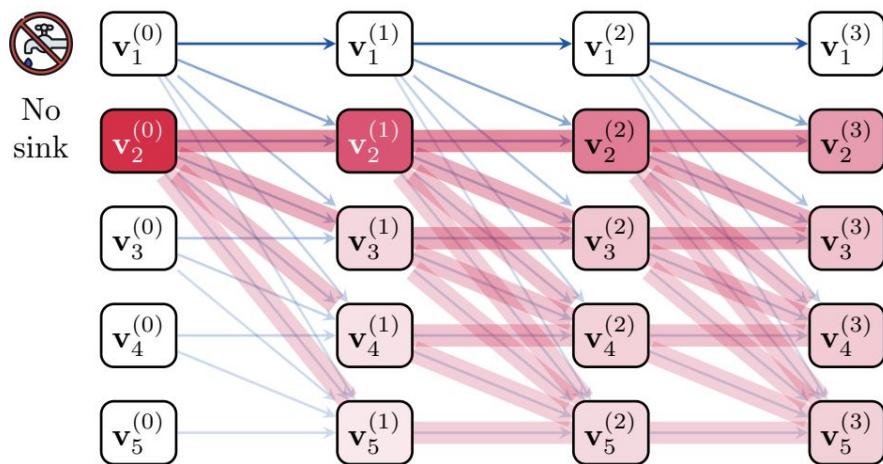
$\text{sim}(\varphi(\mathbf{q}_i), \varphi(\mathbf{k}_j))$	Z_i	$\text{Sink}_1^\epsilon(\%)$	valid loss
$\exp\left(\frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d_h}}\right)$	$\sum_{j'=1}^i \exp\left(\frac{\mathbf{q}_i \mathbf{k}_{j'}^\top}{\sqrt{d_h}}\right)$	18.18	3.73
$\text{sigmoid}\left(\frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d_h}}\right)$	1	0.44*	3.70
$\text{sigmoid}\left(\frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d_h}}\right)$	$\sum_{j'=1}^i \text{sigmoid}\left(\frac{\mathbf{q}_i \mathbf{k}_{j'}^\top}{\sqrt{d_h}}\right)$	30.24	3.74
$\text{elu}\left(\frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d_h}}\right) + 1$	1	0.80*	3.69
$\text{elu}\left(\frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d_h}}\right) + 1$	$\sum_{j'=1}^i \text{elu}\left(\frac{\mathbf{q}_i \mathbf{k}_{j'}^\top}{\sqrt{d_h}}\right) + 1$	-	-
$\frac{(\text{elu}(\mathbf{q}_i) + 1)(\text{elu}(\mathbf{k}_j) + 1)^\top}{\sqrt{d_h}}$	$\sum_{j'=1}^i \frac{(\text{elu}(\mathbf{q}_i) + 1)(\text{elu}(\mathbf{k}_{j'}) + 1)^\top}{\sqrt{d_h}}$	53.65*	4.19
$\frac{(\text{elu}(\mathbf{q}_i) + 1)(\text{elu}(\mathbf{k}_j) + 1)^\top}{\sqrt{d_h}}$	1	-	-
$\mathbf{q}_i \mathbf{k}_j^\top$	$\max\left(\left \sum_{j'=1}^i \mathbf{q}_i \mathbf{k}_{j'}^\top\right _1\right)$		

I am attempting to answer ...

- Mechanism understanding of Attention Sink?
- When Attention Sink Emerges in LLMs?
- Why LLMs need Attention Sink?
- Why GPT-OSS and Qwen3-Next consider Attention Sink in the Model Design?

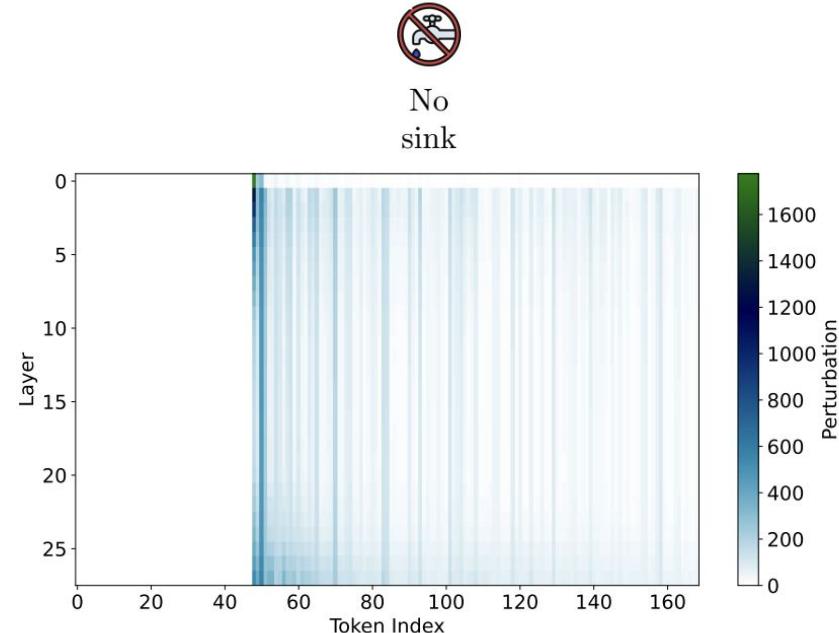
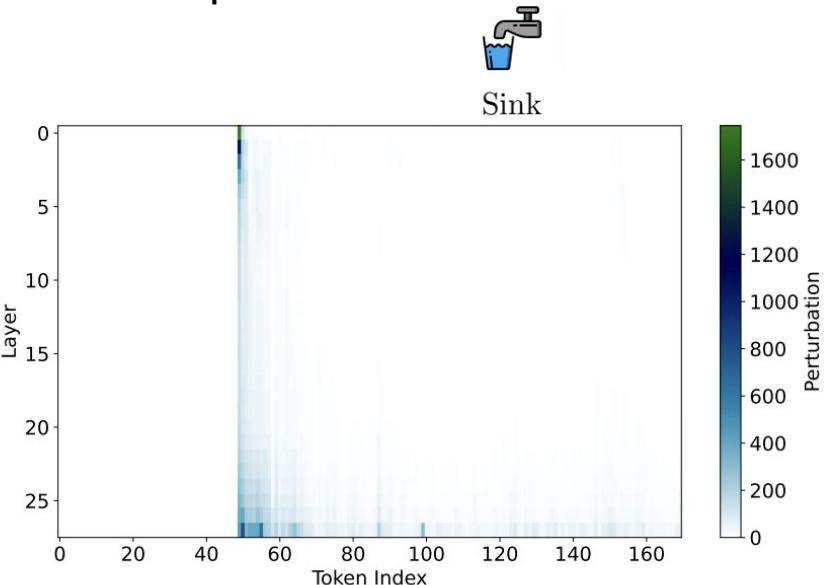
LLMs need attention sink to prevent over-mixing

- Attention blocks try to mix representations
- Attention sink serves as a mechanism to prevent over-mixing (see the paper for theory, longer context needs stronger mechanism)



LLMs need attention sink to prevent over-mixing

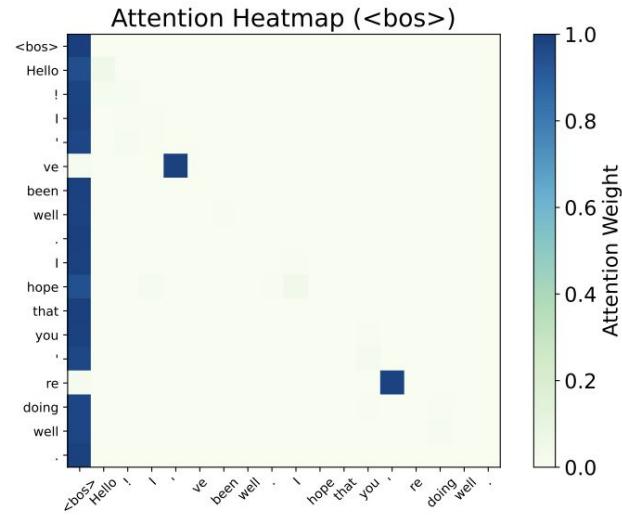
With attention sink, perturbation on one token (“greatest”->“best”) won’t change token representations a lot



Attention sink implements “no-op”

- Attention sink approximates “no-op”: either sharply to attend one important token or attend to the first token
 - From the representation mixing perspective, LLMs need “no-op” to prevent over-mixing

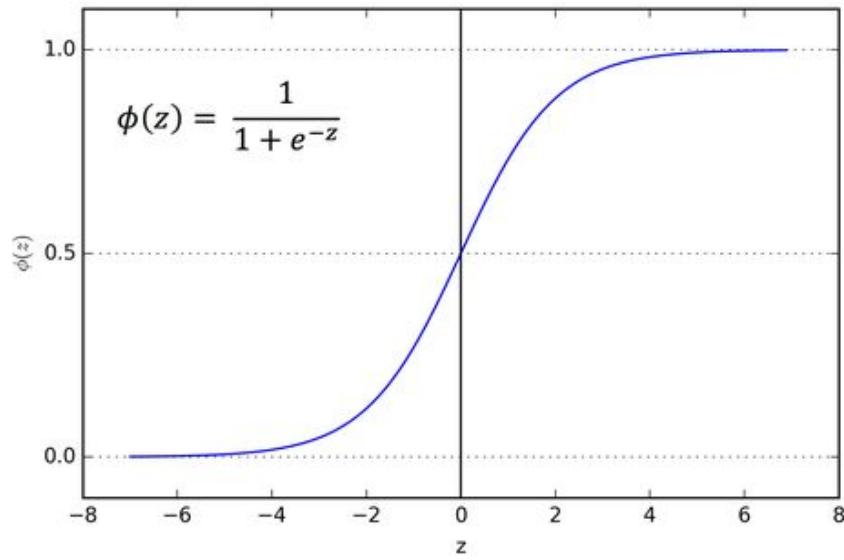
$$\boldsymbol{v}_i^\dagger = \sum_{j=1}^i \alpha_{ij} \boldsymbol{v}_j = \alpha_{i1} \boxed{\boldsymbol{v}_1} + \sum_{j \neq 1}^i \boxed{\alpha_{ij}} \boldsymbol{v}_j$$



Interpreting attention variants using “no-op”

Sigmoid attention allows approximate zero attention

$$\text{Sigmoid} \left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \mathbf{K}^{l,h^\top} + \mathbf{M} \right) \mathbf{V}^{l,h}$$



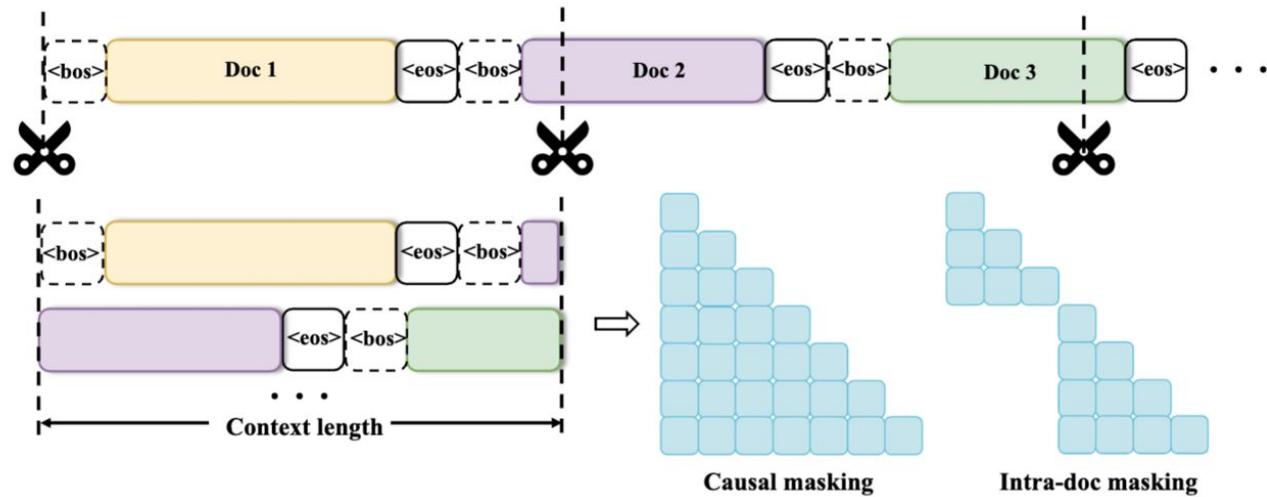
Interpreting attention variants using “no-op”

The following linear attention could have all zero attention scores

$\text{sim}(\varphi(\mathbf{q}_i), \varphi(\mathbf{k}_j))$	Z_i	$\text{Sink}_1^\epsilon(\%)$	valid loss
$\exp\left(\frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d_h}}\right)$	$\sum_{j'=1}^i \exp\left(\frac{\mathbf{q}_i \mathbf{k}_{j'}^\top}{\sqrt{d_h}}\right)$	18.18	3.73
$\frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d_h}}$	$\max\left(\left \sum_{j'=1}^i \frac{\mathbf{q}_i \mathbf{k}_{j'}^\top}{\sqrt{d_h}}\right , 1\right)$	-	-
$\frac{\mathbf{q}_i \mathbf{k}_j^\top}{\sqrt{d_h}}$	1	0.00*	3.99
$\frac{\text{mlp}(\mathbf{q}_i)\text{mlp}(\mathbf{k}_j)^\top}{\sqrt{d_h}}$	$\max\left(\left \sum_{j'=1}^i \frac{\text{mlp}(\mathbf{q}_i)\text{mlp}(\mathbf{k}_{j'})^\top}{\sqrt{d_h}}\right , 1\right)$	0.19*	3.85
$\frac{\text{mlp}(\mathbf{q}_i)\text{mlp}(\mathbf{k}_j)^\top}{\sqrt{d_h}}$	1	0.74*	3.91

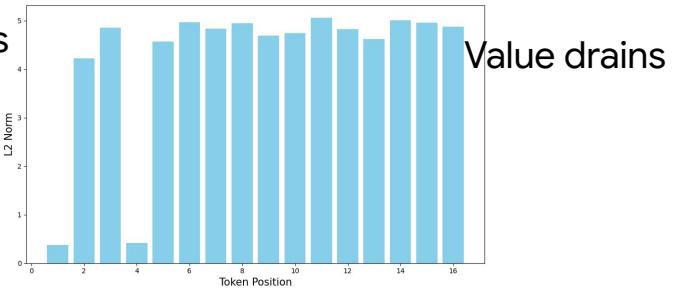
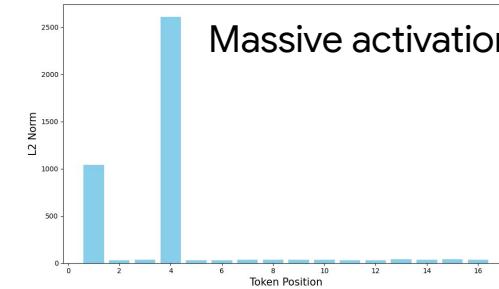
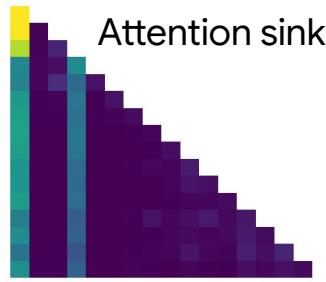
When Attention Sink Attaches to <BOS>

Data packing (fixed <BOS> in the first position will have similar behavior as Gemma)

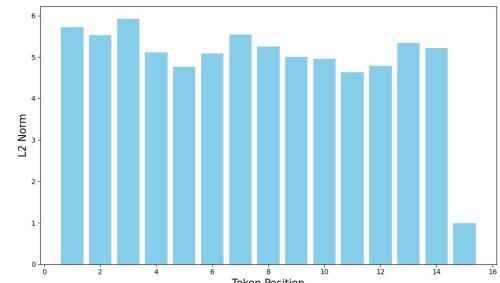
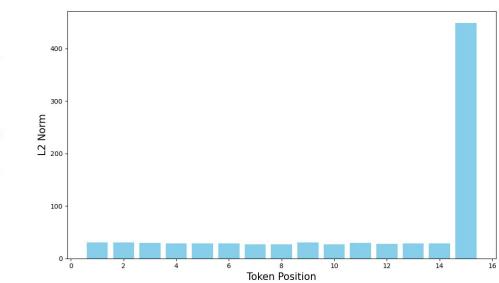
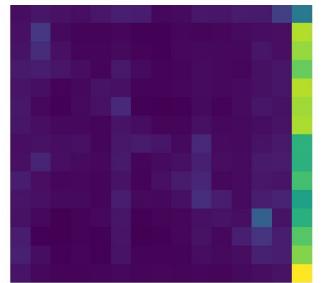


Attention sink / “No-op” widely exists in Transformer family

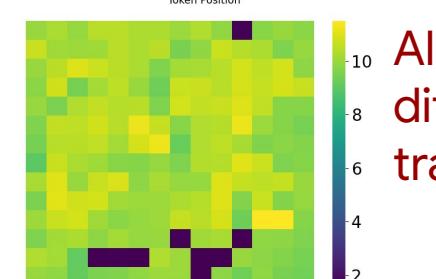
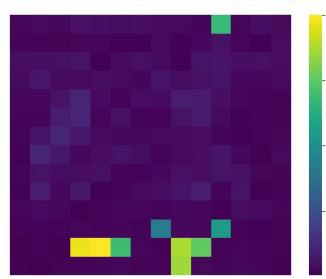
LLaMA



BERT



ViT



Also appear in
diffusion
transformers

I am attempting to answer ...

- Mechanism understanding of Attention Sink?
- When Attention Sink Emerges in LLMs?
- Why LLMs need Attention Sink?
- Why GPT-OSS and Qwen3-Next consider Attention Sink in the Model Design?

GPT-OSS adopts Attention Biases

- Learnable key biases, zero value biases

$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \begin{bmatrix} \mathbf{k}^{*l,h^\top} & \mathbf{K}^{l,h^\top} \end{bmatrix} + \mathbf{M} \right) \begin{bmatrix} \mathbf{0} \\ \mathbf{V}^{l,h} \end{bmatrix}$$

- Softmax off-by-one

$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} \begin{bmatrix} \mathbf{0}^{*l,h^\top} & \mathbf{Q}^{l,h} \mathbf{K}^{l,h^\top} \end{bmatrix} + \mathbf{M} \right) \begin{bmatrix} \mathbf{0} \\ \mathbf{V}^{l,h} \end{bmatrix}$$

- Learnable attention score biases (single number for each head, layer)

$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} \begin{bmatrix} \mathbf{b}^{*l,h^\top} & \mathbf{Q}^{l,h} \mathbf{K}^{l,h^\top} \end{bmatrix} + \mathbf{M} \right) \begin{bmatrix} \mathbf{0} \\ \mathbf{V}^{l,h} \end{bmatrix}$$

$$\mathbf{b}^{*l,h} = b^{*l,h}[1, 1, 1, \dots, 1]$$

GPT-OSS adopts Attention Biases

The first token does not develop strong attention sink, thus mitigating massive activations/outliers

Benefits 1: facilitate quantization, pre-training stability

Pinned

Xiangming Gu ✅ @gu_xiangming · Aug 6

I noticed that @OpenAI added learnable bias to attention logits before softmax. After softmax, they deleted the bias. This is similar to what I have done in my ICLR2025 paper: openreview.net/forum?id=78Nn4...

I used learnable key bias and set corresponding value bias zero. In this way, [Show more](#)

Block	I	J	H	K	V
1	0.1	10.0	0.1	0.1	0.1
2	0.1	4.0	0.1	0.1	0.1
3	0.1	2.0	0.1	0.1	0.1
4	0.1	1.0	0.1	0.1	0.1
5	0.1	0.5	0.1	0.1	0.1
6	0.1	0.3	0.1	0.1	0.1
7	0.1	0.2	0.1	0.1	0.1
8	0.1	0.15	0.1	0.1	0.1
9	0.1	0.1	0.1	0.1	0.1

Attention biases

OpenAI 🤖 @OpenAI · Aug 6

Our open models are here.
Both of them.

openai.com/open-models

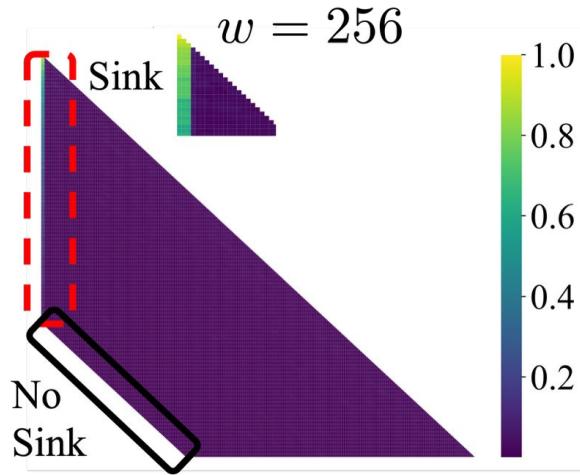
22 185 1.7K 276K

GPT-OSS adopts Attention Biases

- Attention sink only happens in **absolute** first token, not **relative** first token
- Tokens beyond window size have no sinks to attend, possible over-mixing

$$\text{Softmax} \left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \begin{bmatrix} \mathbf{k}^{*,l,h^\top} & \mathbf{K}^{l,h^\top} \end{bmatrix} + \mathbf{M} \right) \begin{bmatrix} \mathbf{0} \\ \mathbf{V}^{l,h} \end{bmatrix}$$

- Facilitate long context, especially in LLMs with alternative shifted window / full attention



Qwen3-Next adopts Gated Attention

$$\text{Sigmoid}(\mathbf{G}^{l,h}) \odot \left[\text{Softmax} \left(\frac{1}{\sqrt{d_h}} \mathbf{Q}^{l,h} \mathbf{K}^{l,h \top} + \mathbf{M} \right) \mathbf{V}^{l,h} \right]$$

Transformations
of inputs

Sigmoid gate allows “no-op”, no need to only rely on attention sink for “no-op” \longrightarrow No attention sink, massive activations, better long context, pre-training stability

Google DeepMind



Thank you for listening!